

MATH 3406 Diary

Introduction

In this diary I will be experimenting with the concepts learned in MATH 3406 and doing the assigned problems. I will be using the [Julia](#) programming language. The diary is maintained through a [Jupyter notebook](#) which lets me combine prose and code in a user-friendly way.

Julia is similar to MATLAB in some regards. It draws inspiration from dynamic languages like Lisp, Perl, Python, and R. Under the hood, Julia operates on the multiple-dispatch paradigm and supports optional typing.

I decided to use this language because I thought it would be interesting to try it out!

Basic operations

```
In [257]: using LinearAlgebra
using Random
using BenchmarkTools # For benchmarking, alternative to tic/ tac; in MATLAB

In [190]: A = [1 2 3; 4 1 6; 7 8 1]

Out [190]: 3×3 Array{Float64,2}:
1 2 3
4 1 6
7 8 1

In [41]: tr(A)

Out [41]: 3

In [51]: det(A)

Out [51]: 104.0

In [61]: inv(A)

Out [61]: 3×3 Array{Float64,2}:
-0.451923  0.211538  0.0865895
0.365385 -0.192308  0.057623
0.144058 -0.057623  0.0673077

In [191]: A * inv(A)

Out [191]: 3×3 Array{Float64,2}:
1.0 0.0 -2.77556e-17
0.0 1.0 0.0
-3.88578e-16 -1.38777e-17 1.0

Note how this is not exactly I3 but is very close to it

In [184]: rref(A)

Out [184]: 2×2 Array{Float64,2}:
0.0 1.0
0.0 0.0
```

Solving a linear system

```
In [187]: A = [1 2 3; 4 1 6; 7 8 1]
b = [5; -2; 9]
A\b

Out [187]: 3-element Array{Float64,1}:
1.0
1.0
2.0

Lets try it out with a singular matrix

In [168]: A = [1 2 3; 4 5 6; 7 8 9]
b = [5; -2; 9]
A\b

SingularException(3)

Stacktrace:
 [1] checknonsingular at ./Users/julia/buildbot/worker/package_macos64/build/usr/share/julia/stdlib/v1.5/LinearAlgebra/src/factorization.jl:119 [inlined]
 [2] checknonsingular at ./Users/julia/buildbot/worker/package_macos64/build/usr/share/julia/stdlib/v1.5/LinearAlgebra/src/factorization.jl:121 [inlined]
 [3] \{::Array{Float64,2}, ::Val{true}; check::Bool) at ./Users/julia/buildbot/worker/package_macos64/build/usr/share/julia/stdlib/v1.5/LinearAlgebra/src/lu.jl:185
 [4] #lu#136 at ./Users/julia/buildbot/worker/package_macos64/build/usr/share/julia/stdlib/v1.5/LinearAlgebra/src/lu.jl:272 [inlined] (repeats 2 times)
 [5] lu at ./Users/julia/buildbot/worker/package_macos64/build/usr/share/julia/stdlib/v1.5/LinearAlgebra/src/lu.jl:272 [inlined] (repeats 2 times)
 [6] \{::Array{Float64,2}, ::Array{Float64,1}) at ./Users/julia/buildbot/worker/package_macos64/build/usr/share/julia/stdlib/v1.5/LinearAlgebra/src/generic.jl:1116
 [7] top-level scope at ./In[188]:3
 [8] include_string(::Function, ::Module, ::String, ::String) at ./loading.jl:1091

Note how Julia throws a SingularException. This is because it checks if the matrix is singular before attempting to perform the calculation. Under the hood, Julia uses an LU decomposition for non-triangular square matrices (which it fails to compute because  $\hat{A}$  is singular).
```

LU Decomposition

```
In [192]: # Tri-diagonal matrix
A = [1 0 0;
      -1 2 1;
      0 0 -2;
      0 0 -2;
      0 0 -2]
P = lu(A)

Out [192]: LU{Float64,Array{Float64,2}}
L factor:
4×4 Array{Float64,2}:
1.0 0.0 0.0 0.0
-1.0 1.0 0.0 0.0
0.0 -1.0 1.0 0.0
0.0 0.0 -1.0 1.0
U factor:
4×4 Array{Float64,2}:
1.0 -1.0 0.0 0.0
0.0 1.0 -1.0 0.0
0.0 0.0 1.0 -1.0
0.0 0.0 0.0 1.0

In [148]: A[P,P,:] == F.L * F.U

Out [148]: true

The above decomposition did not require any row-exchanges.
```

Note how in the above case, we needed to permute the original matrix A . The permutation specifically is a row-exchange between R_1 and R_2

```
In [151]: A = [0 0 1; 0 2 0; 3 0 1]
P = lu(A)

SingularException(1)

Stacktrace:
 [1] checknonsingular at ./Users/julia/buildbot/worker/package_macos64/build/usr/share/julia/stdlib/v1.5/LinearAlgebra/src/factorization.jl:119 [inlined]
 [2] checknonsingular at ./Users/julia/buildbot/worker/package_macos64/build/usr/share/julia/stdlib/v1.5/LinearAlgebra/src/factorization.jl:121 [inlined]
 [3] lu!{::Array{Float64,2}, ::Val{true}; check::Bool) at ./Users/julia/buildbot/worker/package_macos64/build/usr/share/julia/stdlib/v1.5/LinearAlgebra/src/lu.jl:185
 [4] #lu#136 at ./Users/julia/buildbot/worker/package_macos64/build/usr/share/julia/stdlib/v1.5/LinearAlgebra/src/lu.jl:273 [inlined]
 [5] lu at ./Users/julia/buildbot/worker/package_macos64/build/usr/share/julia/stdlib/v1.5/LinearAlgebra/src/lu.jl:273 [inlined]
 [6] \{::Array{Float64,2}, ::Diagonal{Bool,Array{Float64,1}}) at ./Users/julia/buildbot/worker/package_macos64/build/usr/share/julia/stdlib/v1.5/LinearAlgebra/src/diagonal.jl:192
 [7] top-level scope at ./In[151]:3
 [8] include_string(::Function, ::Module, ::String, ::String) at ./loading.jl:1091

Here we tried to perform an LU decomposition on a singular matrix. Note how Julia raises a SingularException. This is because by default Julia checks if the matrix is singular before performing a decomposition.
```

However, we know that every matrix admits an *LUP* decomposition, where P is the permutation matrix. We can disable the checking by passing the `check=false` parameter

```
In [156]: F = lu(A, check=false)
display(F.L)
display(F.P)
display(F.U)

3×3 Array{Float64,2}:
1.0 0.0 0.0
0.0 1.0 0.0
0.0 0.0 1.0
3×3 Array{Float64,2}:
0.0 0.0 1.0
0.0 1.0 0.0
0.0 0.0 2.0
3×3 Array{Float64,2}:
1.0 0.0 0.0
0.0 0.0 1.0
0.0 1.0 0.0
Failed factorization of type LU{Float64,Array{Float64,2}}

Here we see the  $L$ ,  $U$ , and  $P$  factors. It even prints a message saying that the factorization failed (which is expected for this example)
```

Uniqueness of LU factorization

If the matrix A is square and invertible, then there exists a unique *LU* decomposition provided we restrict L to be uni-triangular (i.e., 1s on the diagonal).

However, if A is not invertible, we cannot make such a guarantee. Here is an example.

```
In [167]: A = [0 1; 0 2]

Out [167]: 2×2 Array{Float64,2}:
0 1
0 2

A is not singular. Here are two possible LU decompositions:

In [170]: L1, U1 = lu(A, check=false)
display(L1)
display(U1)

2×2 Array{Float64,2}:
1.0 0.0
0.0 1.0
2×2 Array{Float64,2}:
0 1
0 2

In [179]: L2 = [1 0; 0 1]
U2 = [0 1; 0 1]
display(L2)
display(U2)

2×2 Array{Float64,2}:
1 0
0 1
2×2 Array{Float64,2}:
0 1
0 1

In [181]: L1 * U1 == A

Out [181]: true

In [182]: L2 * U2 == A

Out [182]: true
```

HW 1

Section 1.3

30

For this question, I modeled the system as $\hat{A}x = b$ and used the left division operator `\` to solve it

```
In [111]: # part 1
A = [1 1 1; 1 2 2; 2 3 4]
b = [6; 16; 1]
x = A\b

Out [111]: 3-element Array{Float64,1}:
1.0
3.0
2.0

In [112]: # part 2
A = [1 1 1; 1 2 2; 2 3 4]
b = [7; 14; 3]
x = A\b

Out [112]: 3-element Array{Float64,1}:
2.0
4.0
2.0
```

32

```
In [204]: L,U,p = lu(rand(3,3))

Out [204]: LU{Float64,Array{Float64,2}}
3×3 Array{Float64,2}:
0.3138241 1.0 0.0
0.614814 0.778507 1.0
U factor:
3×3 Array{Float64,2}:
0.779435 0.5035 0.403196
0.0 0.246302 0.88074
0.0 0.0 0.601532

In [205]: U[p[1,1]]

Out [205]: 0.7794346449731999
```

```
In [134]: res = [0, 0, 0]
n = 100000
for i = 1:n
    A = rand(3,3)
    L,U,p = lu(A)
    res += abs.(diag(U)[p])
end
res/n

Out [134]: 3-element Array{Float64,1}:
0.53835456101028
0.5404040696892417
0.539553302227193
```

The average values for each pivot are about 0.5

Section 1.4

21

```
In [198]: A = [0.5 0.5; 0.5 0.5]
println("A^2 = ", A^2)
println("A^3 = ", A^3)

A^2 = [0.5 0.5; 0.5 0.5]
A^3 = [0.5 0.5; 0.5 0.5]

In [199]: B = [1 0; 0 -1]
println("B^2 = ", B^2)
println("B^3 = ", B^3)

B^2 = [1 0; 0 1]
B^3 = [1 0; 0 -1]
```

```
In [200]: C = A*B
println("C = ", C)
println("C^2 = ", C^2)
println("C^3 = ", C^3)

C = [0.5 -0.5; 0.5 -0.5]
C^2 = [0.0 0.0; 0.0 0.0]
C^3 = [0.0 0.0; 0.0 0.0]

B^k = [1 0; 0 (-1)^k]

C^k = [0.5 -0.5; 0.5 -0.5] if k = 1
      [0 0] otherwise

k ∈ N
```

59

```
In [58]: A = I(3)

Out [58]: 3×3 Diagonal{Bool,Array{Bool,1}}:
1 0 0
0 1 0
0 0 1

In [63]: v = [3;4;5]

Out [63]: 3-element Array{Float64,1}:
3
4
5

In [65]: v*

Out [65]: 1×3 Adjoint{Int64,Array{Int64,1}}:
3 4 5
```

```
In [66]: A*v

Out [66]: 3-element Array{Float64,1}:
3
4
5

In [67]: v * v

Out [67]: 50
```

```
In [68]: v*A

Stacktrace:
 [1] check_broadcast_shape at ./broadcast.jl:520 [inlined]
 [2] check_broadcast_shape at ./broadcast.jl:521 [inlined]
 [3] check_broadcast_axes at ./broadcast.jl:523 [inlined]
 [4] instantiate at ./broadcast.jl:269 [inlined]
 [5] materialize! at ./broadcast.jl:285 [inlined]
 [6] materialize! at ./broadcast.jl:285 [inlined]
 [7] rmul!{::Array{Float64,2}, ::Diagonal{Bool,Array{Bool,1}}) at ./Users/julia/buildbot/worker/package_macos64/build/usr/share/julia/stdlib/v1.5/LinearAlgebra/src/diagonal.jl:192
 [8] include_string(::Function, ::Module, ::String, ::String) at ./loading.jl:1091
```

60

```
In [69]: A = ones(4,4)

Out [69]: 4×4 Array{Float64,2}:
1.0 1.0 1.0 1.0
1.0 1.0 1.0 1.0
1.0 1.0 1.0 1.0
1.0 1.0 1.0 1.0

In [70]: v = ones(4,1)

Out [70]: 4×1 Array{Float64,2}:
1.0
1.0
1.0
1.0
```

```
In [71]: A * v

Out [71]: 4×1 Array{Float64,2}:
4.0
4.0
4.0
4.0

In [73]: B = I(4) * ones(4,4)

Out [73]: 4×4 Array{Float64,2}:
2.0 1.0 1.0 1.0
1.0 2.0 1.0 1.0
1.0 1.0 2.0 1.0
1.0 1.0 1.0 2.0

In [74]: w = zeros(4,1) + 2 * ones(4,1)

Out [74]: 4×1 Array{Float64,2}:
2.0
2.0
2.0
2.0
```

```
In [75]: B * w

Out [75]: 4×1 Array{Float64,2}:
4.0
4.0
4.0
4.0

In [76]: A * v

Out [76]: 4×1 Array{Float64,2}:
4.0
4.0
4.0
4.0
```

Section 1.6

12

(a) True

In both cases, we are calculating the inverse of a matrix A using the Gauss-Jordan method.

Case 1: A is upper triangular

Here, A is already in the row-echelon form. We do not need to perform forward elimination. During the first round of backward elimination, we will do a row transformation that subtracts a multiple of the last row from the penultimate row, then last row, the only non-zero entry is the value on the diagonal on both the LHS (if this entry were zero, then A would not be invertible) and RHS (which is the identity matrix). This means that when performing the row transformation, the entries below the diagonal still remain zero, since subtracting zero from a number gives back the number. WLOG, we can apply the same reasoning to conclude that the entries below the diagonal remain zero throughout the entire backward elimination process on both the LHS and RHS. This means that the RHS (which is A^{-1}) is upper-triangular.

$\therefore A^{-1}$ is upper-triangular.

Case 2: A is lower triangular

Here, we only need to perform forward elimination, since the entries above the diagonal are zero. WLOG, we can apply the same reasoning as in the previous case on backward elimination to conclude that the entries above the diagonal remain zero throughout the forward elimination process on both the LHS and RHS.

$\therefore A^{-1}$ is lower-triangular.

From both these cases, we conclude that A^{-1} is triangular.

(b) True

If A is symmetric then,

$$A^T = A$$

But,

$$(A^T)^{-1} = (A^{-1})^T$$

$$\Rightarrow A^{-1} = (A^{-1})^T$$

$\therefore A^{-1}$ is symmetric.

(c) False

Example:

$$A = \begin{bmatrix} 1 & 1 & 0 & 0 \\ 1 & 1 & 1 & 0 \\ 0 & 1 & 1 & 1 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

$$A^{-1} = \begin{bmatrix} 1 & 0 & -1 & 1 \\ 0 & 0 & 1 & -1 \\ -1 & 1 & 0 & 0 \\ 1 & -1 & 0 & 1 \end{bmatrix}$$

(d) False

Example:

$$A = \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

$$A^{-1} = \begin{bmatrix} 1 & 0 \\ 0 & \frac{1}{2} \end{bmatrix}$$

(e) True

The entries of A are given to be fractions. In other words, the entries of A are rational (since every rational number can be expressed in the form $\frac{p}{q}$ where q is non-zero). We know that the set of rational numbers are closed under addition, multiplication, and subtraction. The set of row-transformations (as explained below) we need to apply to find the inverse using the Gauss-Jordan process precisely use addition, multiplication, and subtraction.

Lets say we want to make the entry $\frac{p}{q}$ on row R_i zero by applying a row transformation that subtracts this row from a row R_j (with $j < i$) whose corresponding column entry is $\frac{p}{q}$ (that is non-zero, it wouldn't make sense to subtract a zero entry). We can write this transformation as

$$R_i \rightarrow R_i - \frac{p}{q} R_j$$

Observe that this transformation takes rational numbers to rational numbers. $\frac{a}{b}$ exists since $\frac{a}{b}$ is non-zero.

The corresponding calculation that takes place on the RHS is also rational, since we start off with the identity matrix.

Row-exchanges do not change the rationality of the entries.

By performing a series of the above-mentioned transformations, we will be able to calculate the inverse of A using the Gauss-Jordan process.

```
In [208]: # Calculation for part (c)
A = [
  [ 1 1 1 0 0;
    1 1 1 1 0;
    0 1 1 1 0;
    0 0 1 1 1;
    0 0 0 1 1 ]
]
inv(A)

Out [208]: 4×4 Array{Float64,2}:
1.0 0.0 -1.0 1.0
-1.0 0.0 1.0 -1.0
-1.0 1.0 0.0 -0.0
1.0 -1.0 0.0 1.0

In [214]: # Calculation for part (d)
A = [
  [ 1 0 0;
    0 1 0;
    0 0 1 ]
]
inv(A)

Out [214]: 2×2 Array{Float64,2}:
1.0 0.0
0.0 0.5
```

32

```
In [76]: A = 5 * I(4) - ones(4,4)

Out [76]: 4×4 Array{Float64,2}:
4.0 -1.0 -1.0 -1.0
-1.0 4.0 -1.0 -1.0
-1.0 -1.0 4.0 -1.0
-1.0 -1.0 -1.0 4.0

In [78]: inv(A)

Out [78]: 4×4 Array{Float64,2}:
0.4 0.2 0.2 0.2
0.2 0.4 0.2 0.2
0.2 0.2 0.4 0.2
0.2 0.2 0.2 0.4
a = 0.4 and b = 0.2
```

```
In [80]: A = 6 * I(5) - ones(5,5)

Out [80]: 5×5 Array{Float64,2}:
5.0 -1.0 -1.0 -1.0 -1.0
-1.0 5.0 -1.0 -1.0 -1.0
-1.0 -1.0 5.0 -1.0 -1.0
-1.0 -1.0 -1.0 5.0 -1.0
-1.0 -1.0 -1.0 -1.0 5.0

In [82]: inv(A)

Out [82]: 5×5 Array{Float64,2}:
0.333333 0.166667 0.166667 0.166667 0.166667
0.166667 0.333333 0.166667 0.166667 0.166667
0.166667 0.166667 0.333333 0.166667 0.166667
0.166667 0.166667 0.166667 0.333333 0.166667
0.166667 0.166667 0.166667 0.166667 0.333333
a = 0.3333 and b = 0.1667
```

47

```
In [84]: A = ones(4,4)
b = rand(4,1)
x = A\b

Warning: Matrix is singular to working precision.

x =
 NaN
 NaN
 NaN
 NaN
Inf

b = ones(4,1)

Out [84]: 4×1 Array{Float64,2}:
1.0
1.0
1.0
1.0

Warning: Matrix is singular to working precision.

x =
 NaN
 NaN
 NaN
 NaN
```

Instead of throwing an exception, MATLAB issues a warning.

69

```
In [265]: A = rand(500, 500)
@btime inv(A)

500×500 Array{Float64,2}:
-0.332058 0.500095 0.198305 0.603398 ... -0.41236 0.868234 -0.158316
-0.273561 -0.237709 0.77494 ... -0.145578 0.738137 -0.14288
-0.277316 -0.545251 0.300348 ... -0.68765 1.20534 -0.320808
-0.0616053 -0.119039 0.329892 ... 0.238535 0.053979 0.112729
-0.0673384 -0.087837 -0.635576 ... -0.796961 2.39993 -1.87481
0.0721308 -0.0716083 0.0177225 ... -0.496813 0.484021 -0.203767
-0.156267 -0.523618 -0.269109 ... -1.62106 1.71222 -0.836199
-0.149155 0.0327229 -0.160138 ... -0.95324 2.68777 0.586257
-0.174877 0.164628 -0.879089 ... 1.73179 -1.10885 -0.962218
0.0773281 0.0235047 -0.311736 ... -0.363162 0.264814 -0.398866
-0.0788951 0.0328693 0.166038 ... -0.228223 0.859169 0.68893
-0.19437 -0.1113074 -0.052644 ... 0.0885218 0.132097 -0.136744
0.021343 -0.024486 0.110858 ... -0.37461 -1.08097 0.881924
-0.189136 -0.185462 0.382118 ... -0.288165 0.110801 -0.192848
0.0630003 -0.107217 -0.211561 ... -0.610211 0.662377 -0.360026

We are using the @btime macro from the BenchmarkTools package which lets us benchmark our functions easily.

For the above two cases, computing the inverse of a random 1000×1000 matrix took longer than a random 500×500 matrix.

In our case, our random A matrices were invertible. This may not always be the case since rand may output a singular matrix (which is not invertible). However, the chance of this happening is very low.
```

70

```
In [260]: I = I(1000)
A = rand(1000, 1000)
B = triu(A)

Out [260]: 1000×1000 Array{Float64,2}:
0.682626 0.500095 0.198305 0.603398 ... 0.901597 0.239783 0.147324
-0.332058 0.500095 0.198305 0.603398 ... -0.41236 0.868234 -0.158316
-0.273561 -0.237709 0.77494 ... -0.145578 0.738137 -0.14288
-0.277316 -0.545251 0.300348 ... -0.68765 1.20534 -0.320808
-0.0616053 -0.119039 0.329892 ... 0.238535 0.053979 0.112729
-0.0673384 -0.087837 -0.635576 ... -0.796961 2.39993 -1.87481
0.0721308 -0.0716083 0.0177225 ... -0.496813 0.484021 -0.203767
-0.156267 -0.523618 -0.269109 ... -1.62106 1.71222 -0.836199
-0.149155 0.0327229 -0.160138 ... -0.95324 2.68777 0.586257
-0.174877 0.164628 -0.879089 ... 1.73179 -1.10885 -0.962218
0.0773281 0.0235047 -0.311736 ... -0.363162 0.264814 -0.398866
-0.0788951 0.0328693 0.166038 ... -0.228223 0.859169 0.68893
-0.19437 -0.1113074 -0.052644 ... 0.0885218 0.132097 -0.136744
0.021343 -0.024486 0.110858 ... -0.37461 -1.08097 0.881924
-0.189136 -0.185462 0.382118 ... -0.288165 0.110801 -0.192848
0.0630003 -0.107217 -0.211561 ... -0.610211 0.662377 -0.360026
```



```
0.0      0.0      0.0      0.0      0.816063  0.571443  0.106854
0.0      0.0      0.0      0.0      0.0      0.572444  0.233029
0.0      0.0      0.0      0.0      0.0      0.0      0.863416
```

```
In [261.]
@btime inv(B)

14.816 ms (2 allocations: 7.63 MiB)
Out[261.]
1000x1000 Array{Float64,2}:
 1.49643 -2.14392 -0.0914981  0.698251 ...  5.40099e198  -2.33354e198
 0.0      2.86482 -0.73491  -2.18832   -1.62421e198  7.01759e197
 0.0      0.0      1.545    -0.804926  1.40366e198  -6.0647e197
 0.0      0.0      0.0      1.30489   -2.01067e199  8.68733e198
 0.0      0.0      0.0      0.0      -2.59761e198  1.12231e198
 0.0      0.0      0.0      0.0      -6.45524e198  -2.78903e198
 0.0      0.0      0.0      0.0      -1.77415e199  -7.66551e198
 0.0      0.0      0.0      0.0      -2.62962e198  -2.13615e198
 0.0      0.0      0.0      0.0      -1.65743e198  7.16104e197
 0.0      0.0      0.0      0.0      -4.54799e197  -1.96499e197
 0.0      0.0      0.0      0.0      -2.01685e198  8.71394e197
 0.0      0.0      0.0      0.0      -3.64057e197  1.57293e197
 0.0      0.0      0.0      0.0      2.94919e198  -1.27422e198
 ...
 0.0      0.0      0.0      0.0      -20.5555     11.6857
 0.0      0.0      0.0      0.0      7.70743     -2.63851
 0.0      0.0      0.0      0.0      -5.28121     1.88249
 0.0      0.0      0.0      0.0      -8.03518     1.03479
 0.0      0.0      0.0      0.0      -4.22291     1.72643
 0.0      0.0      0.0      0.0      -5.80516     0.129397
 0.0      0.0      0.0      0.0      6.17229     -2.30802
 0.0      0.0      0.0      0.0      -0.375153    0.398597
 0.0      0.0      0.0      0.0      6.17229     -2.30802
 0.0      0.0      0.0      0.0      -0.308206    -0.132299
 0.0      0.0      0.0      0.0      -1.04334     0.129397
 0.0      0.0      0.0      0.0      1.48711     -0.401359
 0.0      0.0      0.0      0.0      0.0          1.15819
```

```
In [262.]
@btime B\I

14.605 ms (2 allocations: 7.63 MiB)
Out[262.]
1000x1000 Array{Float64,2}:
 1.49643 -2.14392  0.0914981  0.698251 ...  5.40099e198  -2.33354e198
 0.0      2.86482 -0.73491  -2.18832   -1.62421e198  7.01759e197
 0.0      0.0      1.545    -0.804926  1.40366e198  -6.0647e197
 0.0      0.0      0.0      1.3049    -2.01067e199  8.68733e198
 0.0      0.0      0.0      0.0      -2.59761e198  1.12231e198
 0.0      0.0      0.0      0.0      -6.45524e198  -2.78903e198
 0.0      0.0      0.0      0.0      -1.77415e199  -7.66551e198
 0.0      0.0      0.0      0.0      -2.62962e198  -2.13615e198
 0.0      0.0      0.0      0.0      -1.65743e198  7.16104e197
 0.0      0.0      0.0      0.0      -4.54799e197  -1.96499e197
 0.0      0.0      0.0      0.0      -2.01685e198  8.71394e197
 0.0      0.0      0.0      0.0      -3.64057e197  1.57293e197
 0.0      0.0      0.0      0.0      2.94919e198  -1.27422e198
 ...
 0.0      0.0      0.0      0.0      -20.5555     11.6857
 0.0      0.0      0.0      0.0      7.70743     -2.63851
 0.0      0.0      0.0      0.0      -5.28121     1.88249
 0.0      0.0      0.0      0.0      -8.03518     1.03479
 0.0      0.0      0.0      0.0      -4.22291     1.72643
 0.0      0.0      0.0      0.0      -5.80516     0.129397
 0.0      0.0      0.0      0.0      6.17229     -2.30802
 0.0      0.0      0.0      0.0      -0.375153    0.398597
 0.0      0.0      0.0      0.0      6.17229     -2.30802
 0.0      0.0      0.0      0.0      -0.308206    -0.132299
 0.0      0.0      0.0      0.0      -1.04334     0.129397
 0.0      0.0      0.0      0.0      1.48711     -0.401359
 0.0      0.0      0.0      0.0      0.0          1.15819
```

```
In [263.]
@btime inv(A)

38.903 ms (5 allocations: 8.13 MiB)
Out[263.]
1000x1000 Array{Float64,2}:
 -0.523667 -0.482268 -0.431259 ...  0.337907 -0.177193
 0.0302209  0.0881947  0.0795895 ...  0.125973  0.136548
 -0.2323 -0.175653 -0.0670776 ...  0.272359 -0.0551474
 -0.221193 -0.241093 -0.33921 ...  0.46704 -0.159117
 -0.17256 -0.218634 -0.166742 ...  0.116869 -0.10111
 0.100683  0.165526 -0.251768 ... -0.0164263  0.146091
 -0.117831 -0.221348 -0.221403 ...  0.0552534 -0.10111
 -0.497778 -0.493654 -0.430099 ...  0.506457 -0.300345
 -0.035493 -0.100699  0.00524921 ... -0.0385063 -0.000566493
 0.187716  0.172202 -0.189148 ... -0.186028 -0.0132704
 -0.0657512  0.0135414 -0.0727692 ... -0.20124 -0.183351
 0.186675  0.237737  0.239209 ... -0.00271675  0.111346
 -0.299828 -0.227547 -0.278998 ...  0.141411 -0.240707
 ...
 0.108108  0.150003  0.218897 ... -0.223295  0.0817466
 0.181362  0.202014  0.318621 ... -0.244351  0.196917
 -0.496035 -0.483929 -0.643866 ... -0.393983 -0.269871
 -0.231752 -0.198725 -0.378463 ...  0.367829 -0.0990437
 -0.491347 -0.535643 -0.607982 ...  0.384223 -0.289614
 0.280651  0.210103  0.365192 ... -0.243575  0.222415
 -0.31649 -0.26873 -0.13673 ...  0.152516 -0.21772
 -0.312868 -0.335007 -0.294355 ...  0.131205 -0.248601
 -0.0514909 -0.0730094  0.00637453 ... -0.00946273 -0.00663144
 0.06681 -0.00347339  0.0198718 ...  0.121122 -0.0265237
 -0.311042 -0.2784 -0.159137 ...  0.179849 -0.226274
 -0.164406 -0.218984 -0.216931 ...  0.234302  0.0177381
```

```
In [264.]
@btime A\I

39.896 ms (5 allocations: 8.13 MiB)
Out[264.]
1000x1000 Array{Float64,2}:
 -0.523667 -0.482268 -0.431259 ...  0.337907 -0.177193
 0.0302209  0.0881947  0.0795895 ...  0.125973  0.136548
 -0.2323 -0.175653 -0.0670776 ...  0.272359 -0.0551474
 -0.221193 -0.241093 -0.33921 ...  0.46704 -0.159117
 -0.17256 -0.218634 -0.166742 ...  0.116869 -0.10111
 0.100683  0.165526 -0.251768 ... -0.0164263  0.146091
 -0.117831 -0.221348 -0.221403 ...  0.0552534 -0.10111
 -0.497778 -0.493654 -0.430099 ...  0.506457 -0.300345
 -0.035493 -0.100699  0.00524921 ... -0.0385063 -0.000566493
 0.187716  0.172202 -0.189148 ... -0.186028 -0.0132704
 -0.0657512  0.0135414 -0.0727692 ... -0.20124 -0.183351
 0.186675  0.237737  0.239209 ... -0.00271675  0.111346
 -0.299828 -0.227547 -0.278998 ...  0.141411 -0.240707
 ...
 0.108108  0.150003  0.218897 ... -0.223295  0.0817466
 0.181362  0.202014  0.318621 ... -0.244351  0.196917
 -0.496035 -0.483929 -0.643866 ... -0.393983 -0.269871
 -0.231752 -0.198725 -0.378463 ...  0.367829 -0.0990437
 -0.491347 -0.535643 -0.607982 ...  0.384223 -0.289614
 0.280651  0.210103  0.365192 ... -0.243575  0.222415
 -0.31649 -0.26873 -0.13673 ...  0.152516 -0.21772
 -0.312868 -0.335007 -0.294355 ...  0.131205 -0.248601
 -0.0514909 -0.0730094  0.00637453 ... -0.00946273 -0.00663144
 0.06681 -0.00347339  0.0198718 ...  0.121122 -0.0265237
 -0.311042 -0.2784 -0.159137 ...  0.179849 -0.226274
 -0.164406 -0.218984 -0.216931 ...  0.234302  0.0177381
```

The running times are almost the same.

`\` performs slightly better than `inv`

71

In Julia, rational numbers are created using the `//` operator. We can display floating point numbers as fractions using the `rationalize` function. In the following code blocks, we broadcast the `rationalize` function to the matrix.

```
In [220.]
L = [1 0 0 0; -0.5 1 0 0; 0 -2/3 1 0; 0 0 -3/4 1]
rationalize.(L)
```

```
Out[220.]
4x4 Array{Rational{Int64},2}:
 1//1  0//1  0//1  0//1
 -1//2 1//1  0//1  0//1
 0//1 -2//3 1//1  0//1
 0//1  0//1 -3//4 1//1
```

```
In [221.]
rationalize.(inv(L))
```

```
Out[221.]
4x4 Array{Rational{Int64},2}:
 1//1  0//1  0//1  0//1
 1//2 1//1  0//1  0//1
 1//3 2//3 1//1  0//1
 1//4 1//2 3//4 1//1
```

Testing the pattern

```
In [217.]
L = I(5) - diagm(1:5)\diagm(-1⇒1:5)
rationalize.(L)
```

```
Out[217.]
5x5 Array{Rational{Int64},2}:
 1//1  0//1  0//1  0//1  0//1
 -1//2 1//1  0//1  0//1  0//1
 0//1 -2//3 1//1  0//1  0//1
 0//1  0//1 -3//4 1//1  0//1
 0//1  0//1  0//1 -4//5 1//1
```

```
In [227.]
rationalize.(inv(L))
```

```
Out[227.]
4x4 Array{Rational{Int64},2}:
 1//1  0//1  0//1  0//1
 1//2 1//1  0//1  0//1
 1//3 2//3 3//4 1//1
 1//4 1//2 3//4 1//1
```

The general formula for L is

$$L = I(x) - \text{diagm}(1:x)\backslash \text{diagm}(-1\Rightarrow 1:(x-1))$$

Let's test it out for the next few matrices in the sequence

```
In [250.]
function q71(x)
    I(x) - diagm(1:x)\diagm(-1⇒1:(x-1))
end
```

Out[250.] q71 (generic function with 1 method)

```
In [234.]
L = q71(6)
rationalize.(L)
```

```
Out[234.]
6x6 Array{Rational{Int64},2}:
 1//1  0//1  0//1  0//1  0//1  0//1
 -1//2 1//1  0//1  0//1  0//1  0//1
 0//1 -2//3 1//1  0//1  0//1  0//1
 0//1  0//1 -3//4 1//1  0//1  0//1
 0//1  0//1  0//1 -4//5 1//1  0//1
 0//1  0//1  0//1  0//1 -5//6 1//1
```

```
In [235.]
rationalize.(inv(L))
```

```
Out[235.]
6x6 Array{Rational{Int64},2}:
 1//1  0//1  0//1  0//1  0//1  0//1
 1//2 1//1  0//1  0//1  0//1  0//1
 1//3 2//3 1//1  0//1  0//1  0//1
 1//4 1//2 3//4 1//1  0//1  0//1
 1//5 2//5 3//5 4//5 1//1  0//1
 1//6 1//3 1//2 2//3 5//6 1//1
```

```
In [244.]
L = q71(7)
rationalize.(L)
```

```
Out[244.]
7x7 Array{Rational{Int64},2}:
 1//1  0//1  0//1  0//1  0//1  0//1  0//1
 -1//2 1//1  0//1  0//1  0//1  0//1  0//1
 0//1 -2//3 1//1  0//1  0//1  0//1  0//1
 0//1  0//1 -3//4 1//1  0//1  0//1  0//1
 0//1  0//1  0//1 -4//5 1//1  0//1  0//1
 0//1  0//1  0//1  0//1 -5//6 1//1  0//1
 0//1  0//1  0//1  0//1  0//1 -6//7 1//1
```

```
In [245.]
rationalize.(inv(L))
```

```
Out[245.]
7x7 Array{Rational{Int64},2}:
 1//1  0//1  0//1  0//1  0//1  0//1  0//1
 1//2 1//1  0//1  0//1  0//1  0//1  0//1
 1//3 2//3 1//1  0//1  0//1  0//1  0//1
 1//4 1//2 3//4 1//1  0//1  0//1  0//1
 1//5 2//5 3//5 4//5 1//1  0//1  0//1
 1//6 1//3 1//2 2//3 5//6 1//1  0//1
 1//7 2//7 428914250225764//1000799917193449 4//7 5//7 6//7 1//1
```

Interesting point to note: Julia fails to correctly output `3/7` in the inverse. Instead, it outputs a number that is very close to it.

```
In [249.]
float(428914250225764//1000799917193449) - (3/7)
```

```
Out[249.] 1.66533453369377348e-16
```